

## CS188 Spring 2019 Section 11: Perceptrons and Logistic Regression

## 1 Perceptron

**The Algorithm** The perceptron algorithm works as follows:

1. Initialize all weights to 0:  $\mathbf{w} = \mathbf{0}$
2. For each training sample, with features  $\mathbf{f}(x)$  and class label  $y^* \in \{-1, 1\}$ , do:
  - (a) Take the dot product,  $s$ , between the sample features and the current weights:  $s = \mathbf{w}^\top \mathbf{f}(x)$
  - (b) Predict a class,  $\hat{y}$  for the sample as follows:  
 $\hat{y} = +1$  if  $s \geq 0$ ,  $\hat{y} = -1$  otherwise.
  - (c) Compare the predicted label  $\hat{y}$  to the true label  $y^*$ :
    - If  $\hat{y} = y^*$ , do nothing
    - Otherwise, if  $\hat{y} \neq y^*$ , then update your weights:  $\mathbf{w} \leftarrow \mathbf{w} + y^* * \mathbf{f}(x)$
3. If you went through every training sample without having to update your weights (all samples predicted correctly), then terminate. If any at least one sample was predicted incorrectly, then repeat step 2

**Updating weights**

Let us now examine and justify the procedure for updating our weights.

Recall that in step 2b above, we assigned our predicted label  $\hat{y}$  to be either 1 or -1. To update the weights, we first check if the predicted label is correct. If it is, i.e.  $\hat{y} = y^*$ , then do nothing—"don't fix what's not broken", as they say. When they are not equal, then update the weight vector as follows:

$$\mathbf{w} \leftarrow \mathbf{w} + y^* * \mathbf{f}(x)$$


where  $y^*$  is the true label, which is either 1 or -1, and  $x$  the training sample which we mis-classified.

One way to look at this is to see it as a balancing act. If our weights, when multiplied by our sample's features, give us a negative score  $s$  when we wanted a positive score (i.e.  $y^* = 1$ ), then our weights are probably too small (for positive-valued features, or too large for negative-valued). Since  $y^* = 1$ , we, according to this update rule, will add the feature values to our weights to try and make them closer to an optimal set of weights. If our product yields a positive score, where we wanted a negative score, then our weights are probably too big, and so we would like to decrease them. To do so, noting that  $y^* = -1$ , we will subtract our mis-classified sample from our weight vector.

## 2 Perceptron

You want to predict if movies will be profitable based on their screenplays. You hire two critics A and B to read a script you have and rate it on a scale of 1 to 4. The critics are not perfect; here are five data points including the critics' scores and the performance of the movie:

#	Movie Name	A	B	Profit?
1	Pellet Power	1	1	-
2	Ghosts!	3	2	+
3	Pac is Bac	2	4	+
4	Not a Pizza	3	4	+
5	Endless Maze	2	3	-



- (a) First, you would like to examine the linear separability of the data. Plot the data on the 2D plane above; label profitable movies with + and non-profitable movies with - and determine if the data are linearly separable.
- (b) Now you decide to use a perceptron to classify your data. Suppose you directly use the scores given above as features, together with a bias feature. That is  $f_0 = 1$ ,  $f_1 = \text{score given by A}$  and  $f_2 = \text{score given by B}$ .

Run one pass through the data with the perceptron algorithm, filling out the table below. Go through the data points in order, e.g. using data point #1 at step 1.

step	Weights	Score	Correct?
1	$[-1, 0, 0]$	$-1 \cdot 1 + 0 \cdot 1 + 0 \cdot 1 = -1$	yes
2			
3			
4			
5			

Final weights:

- (c) Have weights been learned that separate the data?
- (d) More generally, irrespective of the training data, you want to know if your features are powerful enough to allow you to handle a range of scenarios. Circle the scenarios for which a perceptron using the features above can indeed perfectly classify movies which are profitable according to the given rules:
- (a) Your reviewers are awesome: if the total of their scores is more than 8, then the movie will definitely be profitable, and otherwise it won't be.
  - (b) Your reviewers are art critics. Your movie will be profitable if and only if each reviewer gives either a score of 2 or a score of 3.
  - (c) Your reviewers have weird but different tastes. Your movie will be profitable if and only if both reviewers agree.

### Q3. Optimization

We would like to classify some data. We have  $N$  samples, where each sample consists of a feature vector  $\mathbf{x} = \{x_1, \dots, x_k\}$  and a label  $y = \{0, 1\}$ .

We introduce a new type of classifier called logistic regression, which produces predictions as follows:

$$P(Y = 1|X) = h(\mathbf{x}) = s\left(\sum_i w_i x_i\right) = \frac{1}{1 + \exp(-(\sum_i w_i x_i))}$$
$$s(\gamma) = \frac{1}{1 + \exp(-\gamma)}$$

where  $s(\gamma)$  is the logistic function,  $\exp x = e^x$ , and  $\mathbf{w} = \{w_1, \dots, w_k\}$  are the learned weights.

Let's find the weights  $w_j$  for logistic regression using stochastic gradient descent. We would like to minimize the following loss function for each sample:

$$L = -[y \ln h(\mathbf{x}) + (1 - y) \ln(1 - h(\mathbf{x}))]$$

(a) Find  $dL/dw_i$ . Hint:  $s'(\gamma) = s(\gamma)(1 - s(\gamma))$ .

(b) Write the stochastic gradient descent update for  $w_i$ . Our step size is  $\eta$ .