

Self-assessment due: Monday 5/6/2019 at 11:59pm (submit via Gradescope)

For the self assessment, **fill in the self assessment boxes in your original submission** (you can download a PDF copy of your submission from Gradescope – be sure to delete any extra title pages that Gradescope attaches). For each subpart where your original answer was correct, write “correct.” Otherwise, write and explain the correct answer. **Do not leave any boxes empty.** If you did not submit the homework (or skipped some questions) but wish to receive credit for the self-assessment, we ask that you first complete the homework without looking at the solutions, and then perform the self assessment afterwards.

Q1. The OMNIBUS

(a) Search

- (i) [true or false] Uniform-cost search will never expand more nodes than A*-search.
- (ii) [true or false] Depth-first search will always expand more nodes than breadth-first search.
- (iii) [true or false] The heuristic $h(n) = 0$ is admissible for every search problem.
- (iv) [true or false] The heuristic $h(n) = 1$ is admissible for every search problem.
- (v) [true or false] The heuristic $h(n) = c(n)$, where $c(n)$ is the true cheapest cost to get from the node n to a goal state, is admissible for every search problem.

(b) CSPs

- (i) [true or false] The most-constrained variable heuristic provides a way to select the next variable to assign in a backtracking search for solving a CSP.
- (ii) [true or false] By using the most-constrained variable heuristic and the least-constraining value heuristic we can solve every CSP in time linear in the number of variables.

(c) Games

- (i) [true or false] When using alpha-beta pruning, it is possible to get an incorrect value at the root node by choosing a bad ordering when expanding children.
- (ii) [true or false] When using alpha-beta pruning, the computational savings are independent of the order in which children are expanded.
- (iii) [true or false] When using expectimax to compute a policy, re-scaling the values of all the leaf nodes by multiplying them all with 10 can result in a different policy being optimal.

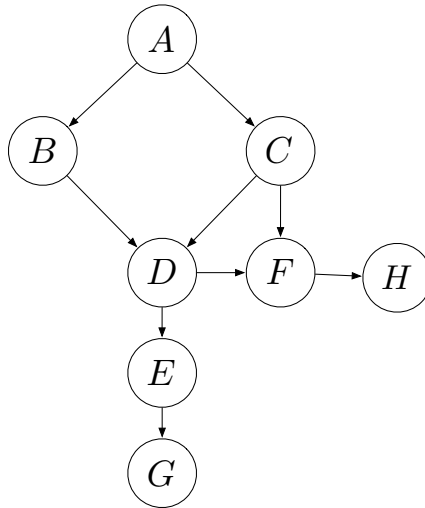
(d) MDPs For this question, assume that the MDP has a finite number of states.

- (i) [true or false] For an MDP (S, A, T, γ, R) if we only change the reward function R the optimal policy is guaranteed to remain the same.
- (ii) [true or false] Value iteration is guaranteed to converge if the discount factor (γ) satisfies $0 < \gamma < 1$.
- (iii) [true or false] Policies found by value iteration are superior to policies found by policy iteration.

(e) Reinforcement Learning

- (i) [true or false] Q-learning can learn the optimal Q-function Q^* without ever executing the optimal policy.
- (ii) [true or false] If an MDP has a transition model T that assigns non-zero probability for all triples $T(s, a, s')$ then Q-learning will fail.

- (f) **Bayes' Nets** For each of the conditional independence assertions given below, circle whether they are guaranteed to be true, guaranteed to be false, or cannot be determined for the given Bayes' net.



$B \perp\!\!\!\perp C$	Guaranteed true	Guaranteed false	Cannot be determined
$B \perp\!\!\!\perp C \mid G$	Guaranteed true	Guaranteed false	Cannot be determined
$B \perp\!\!\!\perp C \mid H$	Guaranteed true	Guaranteed false	Cannot be determined
$A \perp\!\!\!\perp D \mid G$	Guaranteed true	Guaranteed false	Cannot be determined
$A \perp\!\!\!\perp D \mid H$	Guaranteed true	Guaranteed false	Cannot be determined
$B \perp\!\!\!\perp C \mid A, F$	Guaranteed true	Guaranteed false	Cannot be determined
$F \perp\!\!\!\perp B \mid D, A$	Guaranteed true	Guaranteed false	Cannot be determined
$F \perp\!\!\!\perp B \mid D, C$	Guaranteed true	Guaranteed false	Cannot be determined

Q2. Perceptron

- (a) Suppose you have a binary perceptron in 2D with weight vector $\mathbf{w} = r [w_1, w_2]^T$. You are given w_1 and w_2 , and are given that $r > 0$, but otherwise not told what r is. *Assume that ties are broken as positive.*

Can you determine the perceptron's classification of a new example x with known feature vector $f(x)$?

☒ Always ☐ Sometimes ☐ Never

- (b) Now you are learning a multi-class perceptron between 4 classes. The weight vectors are currently $[1, 0]^T$, $[0, 1]^T$, $[-1, 0]^T$, $[0, -1]^T$ for the classes A, B, C, and D. The next training example x has a **label of A** and feature vector $f(x)$.

For the following questions, *do not make any assumptions about tie-breaking.* (Do not write down a solution that creates a tie.)

- (i) Write down a feature vector in which no weight vectors will be updated.

$$f(x) = \begin{bmatrix} 1 \\ 0 \end{bmatrix} \quad \text{Not possible}$$

Any feature vector that points in the direction of w_A more than any other direction, such that $w_A \cdot f(x) > w_i \cdot f(x)$ for $i \neq A$.

- (ii) Write down a feature vector in which **only** \mathbf{w}_A will be updated by the perceptron.

$$f(x) = \begin{bmatrix} \\ \end{bmatrix} \quad \text{Not possible}$$

- (iii) Write down a feature vector in which **only** \mathbf{w}_A and \mathbf{w}_B will be updated by the perceptron.

$$f(x) = \begin{bmatrix} 0 \\ 1 \end{bmatrix} \quad \text{Not possible}$$

Any feature vector that points in the direction of w_B more than any other direction, such that $w_B \cdot f(x) > w_i \cdot f(x)$ for $i \neq B$.

- (iv) Write down a feature vector in which **only** \mathbf{w}_A and \mathbf{w}_C will be updated by the perceptron.

$$f(x) = \begin{bmatrix} -1 \\ 0 \end{bmatrix} \quad \text{Not possible}$$

Any feature vector that points in the direction of w_C more than any other direction, such that $w_C \cdot f(x) > w_i \cdot f(x)$ for $i \neq C$.

The weight vectors are the same as before, but now there is a bias feature with value of 1 for all x and the weight of this bias feature is 0, -2, 1, -1 for classes A, B, C, and D respectively. As before, the next training example x has a **label of A** and a feature vector $f(x)$. The always "1" bias feature is the first entry in $f(x)$.

- (v) Write down a feature vector in which **only** \mathbf{w}_B and \mathbf{w}_C will be updated by the perceptron.

$$f(x) = \begin{bmatrix} 1 \\ \end{bmatrix} \quad \text{Not possible}$$

- (vi) Write down a feature vector in which **only** \mathbf{w}_A and \mathbf{w}_C will be updated by the perceptron.

$$f(x) = \begin{bmatrix} 1 \\ 0 \\ 0 \end{bmatrix} \quad \text{Not possible}$$

Any feature vector that points in the direction of w_C more than any other direction, such that $w_C \cdot f(x) > w_i \cdot f(x)$ for $i \neq C$.